

OpenGL 4

Introduction to Computer Graphics

Andrzej Szymczak

October 12, 2012

Motivation



- ▶ The old way requires sending a lot of commands from the CPU to the GPU
 - ▶ CPU could waste time converting data/commands into the right format
 - ▶ More communication than necessary
 - ▶ Let GPU operate on 'objects' contained in its own memory
 - ▶ Better send large blocks of data using a small number of OpenGL calls (e.g. all vertices at once vs one at a time as in OpenGL 1 `glVertex` function)
- ▶ Since programmable shaders are available:
 - ▶ Why hard code vertex attributes any more? Give the programmer control over what they are
 - ▶ *Require* the programmer provide shaders; after all, the traditional functionality can be replicated this way

Specifying vertices: the new way



- ▶ In many cases, the models you render don't change for a number of frames
- ▶ Keep the vertex data in the GPU's memory to save time (less CPU-GPU communication)
- ▶ Generic attributes
 - ▶ Let the programmer specify any number of attributes of different types for each vertex
 - ▶ Each attribute will have an *index*
 - ▶ Programmer-provided vertex shader will recognize attributes by index and interpret them properly

Specifying vertices: VAOs and VBOs


▶ Vertex Buffer Objects

- ▶ Arrays of vertex properties (coordinates, normals, colors or anything else) stored in the GPU memory
- ▶ Identified by a VBO handle, an unsigned integer (GLuint type)

▶ Vertex Array Objects

- ▶ Think of them as bundles of VBOs, with instructions on how to build vertices from VBOs' contents
- ▶ Identified by a VAO handle, a GLuint type value (note: this generally the case for all OpenGL 'objects')
- ▶ Each VBO in the bundle holds a different vertex attribute
- ▶ Attributes have *indices*, integer identifiers
- ▶ Different attributes may have different types or number of components (e.g. can be 4D vectors, 2D vectors, floats, integers)
- ▶ Generally, when creating a 'vertex', the GPU will assemble its attributes by looking them up from all active VBOs for the current active VAO

VAOs and VBOs



Creating a useful VAO

- ▶ Generate a handle for the VAO
 - ▶ `GLuint vao;`
 - ▶ `glGenVertexArrays(1,&vao);` // first arg is # handles to generate

Creating a useful VAO

- ▶ Generate a handle for the VAO
 - ▶ `GLuint vao;`
 - ▶ `glGenVertexArrays(1,&vao);` // first arg is # handles to generate
- ▶ Bind the VAO; basically, this means that you want to work on/with your VAO now
 - ▶ `glBindVertexArray(vao);`

Creating a useful VAO

- ▶ Generate a handle for the VAO
 - ▶ `GLuint vao;`
 - ▶ `glGenVertexArrays(1,&vao);` // first arg is # handles to generate
- ▶ Bind the VAO; basically, this means that you want to work on/with your VAO now
 - ▶ `glBindVertexArray(vao);`
- ▶ Now, we need to attach some VBOs to our VAO
 - ▶ Here we'll need one for coordinates and one for normals
 - ▶ Start off by creating handles for them
 - ▶ `GLuint vbo[2];`
 - ▶ `glGenBuffers(2,vbo);` //note similarity to `glGenVertexArrays`

Creating a useful VAO

- ▶ Generate a handle for the VAO
 - ▶ `GLuint vao;`
 - ▶ `glGenVertexArrays(1,&vao);` // first arg is # handles to generate
- ▶ Bind the VAO; basically, this means that you want to work on/with your VAO now
 - ▶ `glBindVertexArray(vao);`
- ▶ Now, we need to attach some VBOs to our VAO
 - ▶ Here we'll need one for coordinates and one for normals
 - ▶ Start off by creating handles for them
 - ▶ `GLuint vbo[2];`
 - ▶ `glGenBuffers(2,vbo);` //note similarity to `glGenVertexArrays`
- ▶ Bind the first VBO
 - ▶ `glBindBuffer(GL_ARRAY_BUFFER,vbo[0]);`

Creating a useful VAO

- ▶ Generate a handle for the VAO
 - ▶ `GLuint vao;`
 - ▶ `glGenVertexArrays(1,&vao);` // first arg is # handles to generate
- ▶ Bind the VAO; basically, this means that you want to work on/with your VAO now
 - ▶ `glBindVertexArray(vao);`
- ▶ Now, we need to attach some VBOs to our VAO
 - ▶ Here we'll need one for coordinates and one for normals
 - ▶ Start off by creating handles for them
 - ▶ `GLuint vbo[2];`
 - ▶ `glGenBuffers(2,vbo);` //note similarity to `glGenVertexArrays`
- ▶ Bind the first VBO
 - ▶ `glBindBuffer(GL_ARRAY_BUFFER,vbo[0]);`
- ▶ Send data from an array in the main memory to the GPU memory
 - ▶ `glBufferData(GL_ARRAY_BUFFER,
12*3*3*sizeof(GLfloat),coordarray,GL_STATIC_DRAW);`

Creating a useful VAO, cont.

- ▶ The current VBO (vbo[0]) stores attribute
 - ▶ with index 0
 - ▶ with 3 coordinates per vertex
 - ▶ of GLfloat type
 - ▶ there is no need to normalize it
 - ▶ Here is how to put this info into the VAO:
`glVertexAttribPointer(0,3,GL_FLOAT,GL_FALSE,0,0);`

Creating a useful VAO, cont.

- ▶ The current VBO (vbo[0]) stores attribute
 - ▶ with index 0
 - ▶ with 3 coordinates per vertex
 - ▶ of GLfloat type
 - ▶ there is no need to normalize it
 - ▶ Here is how to put this info into the VAO:
`glVertexAttribPointer(0,3,GL_FLOAT,GL_FALSE,0,0);`
- ▶ Bind the second VBO and send data (normals) to it; this is attribute of index 1
 - ▶ `glBindBuffer(GL_ARRAY_BUFFER,vbo[1]);`
 - ▶ `glBufferData(GL_ARRAY_BUFFER,
12*3*3*sizeof(GLfloat),normalarray,GL_STATIC_DRAW);`
 - ▶ `glVertexAttribPointer(1,3,GL_FLOAT,GL_FALSE,0,0);`

Creating a useful VAO, cont.

- ▶ The current VBO (vbo[0]) stores attribute
 - ▶ with index 0
 - ▶ with 3 coordinates per vertex
 - ▶ of GLfloat type
 - ▶ there is no need to normalize it
 - ▶ Here is how to put this info into the VAO:
`glVertexAttribPointer(0,3,GL_FLOAT,GL_FALSE,0,0);`
- ▶ Bind the second VBO and send data (normals) to it; this is attribute of index 1
 - ▶ `glBindBuffer(GL_ARRAY_BUFFER,vbo[1]);`
 - ▶ `glBufferData(GL_ARRAY_BUFFER,12*3*3*sizeof(GLfloat),normalarray,GL_STATIC_DRAW);`
 - ▶ `glVertexAttribPointer(1,3,GL_FLOAT,GL_FALSE,0,0);`
- ▶ Enable both attributes
 - ▶ `glEnableVertexAttribArray(0);`
 - ▶ `glEnableVertexAttribArray(1);`

Creating a useful VAO, cont.

- ▶ Finally, unbind the VBO and VAO – we are done working with them
 - ▶ `glBindBuffer(GL_ARRAY_BUFFER,0);`
 - ▶ `glBindVertexArray(0);`

Using our VAO

- ▶ Reminder: the goal is to send these vertices into the pipeline
- ▶ First, bind the VAO
 - ▶ `glBindVertexArray(vao);`
- ▶ Send the vertices
 - ▶ `glDrawArrays(GL_TRIANGLES, 0, 12*3);`
 - ▶ First argument: how to interpret the vertex stream (relevant to primitive setup)
 - ▶ Next two: range of indices to use (note that the same index i is used to pick attributes of the vertex from all VBOs)
- ▶ Unbind the VAO when done
 - ▶ `glBindVertexArray(0);`

Generalities

- ▶ OpenGL is full of different kinds of objects
- ▶ They store data used for rendering and/or some state, basically information on how to interpret the data
- ▶ Working with objects of any type generally requires similar steps
 - ▶ Bind
 - ▶ Work
 - ▶ Unbind
- ▶ The sample code does not free memory used by objects!
- ▶ OK since it only creates a small number of objects at startup
 - ▶ In your project, do the same!
 - ▶ Don't create unnecessary objects
 - ▶ Don't keep *recreating* something that already exists!
 - ▶ ... especially for each frame (you'll have a memory leak)
- ▶ In more complex code, you need to delete objects (such as programs, shaders, VBOs, VAOs etc)
 - ▶ See `glDelete*` functions in OpenGL manual pages

Vertex program

- ▶ Capture input vertex attributes; location qualifier provides the index information
 - ▶ `layout(location=0) in vec3 coord;`
 - ▶ `layout(location=1) in vec3 normal;`
- ▶ We'll send just one output attribute with the processed vertex
 - ▶ `flat` qualifier requests flat 'interpolation'; alternatives: `noperspective` or `smooth`
 - ▶ `flat out float NdotL;`
- ▶ Declare uniform variables – they can be set from the C code using `glUniform*`
 - ▶ `uniform mat4 ModelViewMatrix;`
 - ▶ `uniform mat4 ProjectionMatrix;`
 - ▶ `uniform mat3 NormalMatrix;`
 - ▶ `uniform vec3 LightLocation;`
 - ▶ Note uniform variables can't be changed from the shader

Vertex program

- ▶ Main function

- ▶ Goals

- ▶ Compute values of output variables (here: `NdotL`)
 - ▶ Assign coordinates of the projected vertex (in the homogenous coordinates – typically you don't want to divide by the homogenous coordinate here) to the built-in output variable `vec4 gl_Position`
 - ▶ Note that `gl_Position` is used for clipping and rasterization

- ▶ Code:

- ▶

```
void main() {
```
 - ▶

```
    vec4 WorldCoord = ModelViewMatrix *  
    vec4(coord,1.0);
```
 - ▶

```
    vec3 L = normalize(LightLocation -  
    WorldCoord.xyz);
```
 - ▶

```
    vec3 WorldNormal = NormalMatrix*normal;
```
 - ▶

```
    vec3 N = normalize(WorldNormal);
```
 - ▶

```
    NdotL = dot(N,L);
```
 - ▶

```
    gl_Position = ProjectionMatrix*WorldCoord;
```
 - ▶

```
}
```

Fragment Program

- ▶ Capture *interpolated* output of the vertex shader
 - ▶ This is how vertex and fragment shaders communicate!
 - ▶ Names and qualifiers need to match
 - ▶ `flat in float NdotL;`
 - ▶ recall that in the vertex shader we had `flat out float NdotL;`

Fragment Program

- ▶ Capture *interpolated* output of the vertex shader
 - ▶ This is how vertex and fragment shaders communicate!
 - ▶ Names and qualifiers need to match
 - ▶ `flat in float NdotL;`
 - ▶ recall that in the vertex shader we had `flat out float NdotL;`
- ▶ Declare all uniform variables (they can be set from CPU)
 - ▶ `uniform float LightIntensity;`
 - ▶ `uniform float AmbientIntensity;`
 - ▶ `uniform vec3 DiffuseAndAmbientCoefficient;`

Fragment Program

- ▶ Capture *interpolated* output of the vertex shader
 - ▶ This is how vertex and fragment shaders communicate!
 - ▶ Names and qualifiers need to match
 - ▶ `flat in float NdotL;`
 - ▶ recall that in the vertex shader we had `flat out float NdotL;`
- ▶ Declare all uniform variables (they can be set from CPU)
 - ▶ `uniform float LightIntensity;`
 - ▶ `uniform float AmbientIntensity;`
 - ▶ `uniform vec3 DiffuseAndAmbientCoefficient;`
- ▶ Declare the output variables; if there is exactly one of type `vec4`, the components will be treated as the RGBA values for the fragment
 - ▶ `out vec4 fragcolor;`

Fragment Program

- ▶ The main function is very simple:

- ▶

```
void main() {
```
- ▶

```
    fragcolor = vec4( (LightIntensity*
```



```
        (NdotL > 0.0 ?  NdotL : 0.0) + AmbientIntensity) *  
        DiffuseAndAmbientCoefficient, 1);
```
- ▶

```
}
```

Compiling and linking vertex and fragment programs

- ▶ Create handles
 - ▶ `GLuint vid = glCreateShader(GL_VERTEX_SHADER);`
 - ▶ `GLuint fid = glCreateShader(GL_FRAGMENT_SHADER);`
- ▶ Specify source
 - ▶ `glShaderSource(vid,1,&VS,NULL);`
 - ▶ `glShaderSource(fid,1,&FS,NULL);`
 - ▶ This works if VS and FS are C-style null-terminated strings (char* type)
- ▶ Compile shaders
 - ▶ `glCompileShader(vid);`
 - ▶ `glCompileShader(fid);`
- ▶ Create program handle (it will include complete set of shaders)
 - ▶ `GLuint ProgramHandle = glCreateProgram();`
- ▶ Attach our vertex and fragment shaders to the program
 - ▶ `glAttachShader(ProgramHandle,vid);`
 - ▶ `glAttachShader(ProgramHandle,fid);`

Compiling and linking vertex and fragment programs

▶ Link

- ▶ `glLinkProgram(ProgramHandle);`
- ▶ This is where vertex shader outputs are matched with fragment shader outputs and any shared uniform variables are merged

Using the program

- ▶ `glUseProgram(ProgramHandle);`
- ▶ Now, the program will be applied to any vertex stream
- ▶ In general, you may have several programs, e.g. in your project you may want to write one for flat shading and one for Gouraud shading; you can use `glUseProgram` to switch between different programs – just pass the handle of the program you want to use as the argument

Uniform variables

- ▶ Each uniform variable has a *location* in the program
- ▶ Use `glGetUniformLocation` to get locations of uniforms
 - ▶ Example: `GLint NormalMatrixLoc = glGetUniformLocation(ProgramHandle, "NormalMatrix");`
 - ▶ Return value is -1 means variable was not found
 - ▶ Could mean a typo
 - ▶ ... or that the compiler figured out that the uniform variable can be optimized out

Uniform variables

- ▶ Each uniform variable has a *location* in the program
- ▶ Use `glGetUniformLocation` to get locations of uniforms
 - ▶ Example: `GLint NormalMatrixLoc = glGetUniformLocation(ProgramHandle,"NormalMatrix");`
 - ▶ Return value is -1 means variable was not found
 - ▶ Could mean a typo
 - ▶ ... or that the compiler figured out that the uniform variable can be optimized out
- ▶ Set values of uniform variables using `glUniform*`
 - ▶ The program has to be active (call `glUseProgram` before!)
 - ▶ `glUniform1f(LightIntensityLoc,0.9);`
 - ▶ `glUniform3f(DiffuseAndAmbientCoefficientLoc,0.0,1.0,0.0);` (this one is of `vec3` type)
 - ▶ Alternative: `glUniform3fv` – use pointer to 3D array of floats as second argument
 - ▶ For matrix uniforms, pass pointer to an array of entries (here, array of size 16 since the matrix is 4×4)
`glUniformMatrix4fv(ModelViewMatrixLoc,1,GL_FALSE,ptr);`

- ▶ A simple header-only library that provides useful math functions
 - ▶ Vector operations
 - ▶ Matrix operations
 - ▶ Replicates functionality of most useful obsolete OpenGL functions such as transformation calls
 - ▶ Also, replicates most useful deprecated GL utilities (GLU) functions

Using GLM to set matrices

- ▶ Setting the projection matrix

```
glm::mat4 PMat =  
.    glm::perspective(8.0f,1.0f,15.0f,25.0f);  
glUniformMatrix4fv(ProjectionMatrixLoc,1,  
.    GL_FALSE,&PMat[0][0]);
```

- ▶ Setting the modelview matrix

```
glm::mat4 MVMat = glm::translate(glm::mat4(1.0),  
.    glm::vec3(0.0,0.0,-20.0)) *  
.    glm::rotate(glm::mat4(1.0),angle1,  
.    glm::vec3(1.0,2.0,3.0)) *  
.    glm::rotate(glm::mat4(1.0),angle2,  
.    glm::vec3(-2.0,-1.0,0.0));  
glUniformMatrix4fv(ModelViewMatrixLoc,1,GL_FALSE,  
.    &MVMat[0][0]);
```

Using GLM to set matrices

- ▶ Setting the normal matrix
 - ▶ In this case normals are changed only by the rotational component – thus, we can use the 3×3 linear part of the modelview matrix as the normal matrix
 - ▶ In general, you would take inverse and transpose of NMat used below before sending it to the uniform variable

```
glm::mat3 NMat(MVMat);  
glUniformMatrix3fv(NormalMatrixLoc,1,  
    GL_FALSE,&NMat[0][0]);
```